


Quality vs. speed trade-off

circled = boolean path of the production model (unified SVM class)

■ API LLM ● Embedding + SVM ▲ Local LLM - - - Pareto frontier



Local-model timings measured on a single consumer GPU (RTX 5070 Ti 16 GB); large local models were partially memory-constrained.